

Ensemble Methods

Principle of ensemble classifier

- Ask experts for their opinion and choose the option with majority vote
- Let's say we have a set of m experts:

$$H = \{f_1, f_2, \dots, f_m\} \quad f_i(x) \in \{P, N\}$$

- The majority vote decision will be:

$$F(x) = \text{sign} \left(\frac{1}{m} \sum_{i=1}^m f_i(x) \right)$$

- Diversity of the expert a key point of this approach

Why to use the ensemble methods

Better performance

Assume that $\forall i, p(f_i(x) \neq y) \leq \mu < 0.5$

Classifiers are independent then the probability of wrong classification by the ensemble

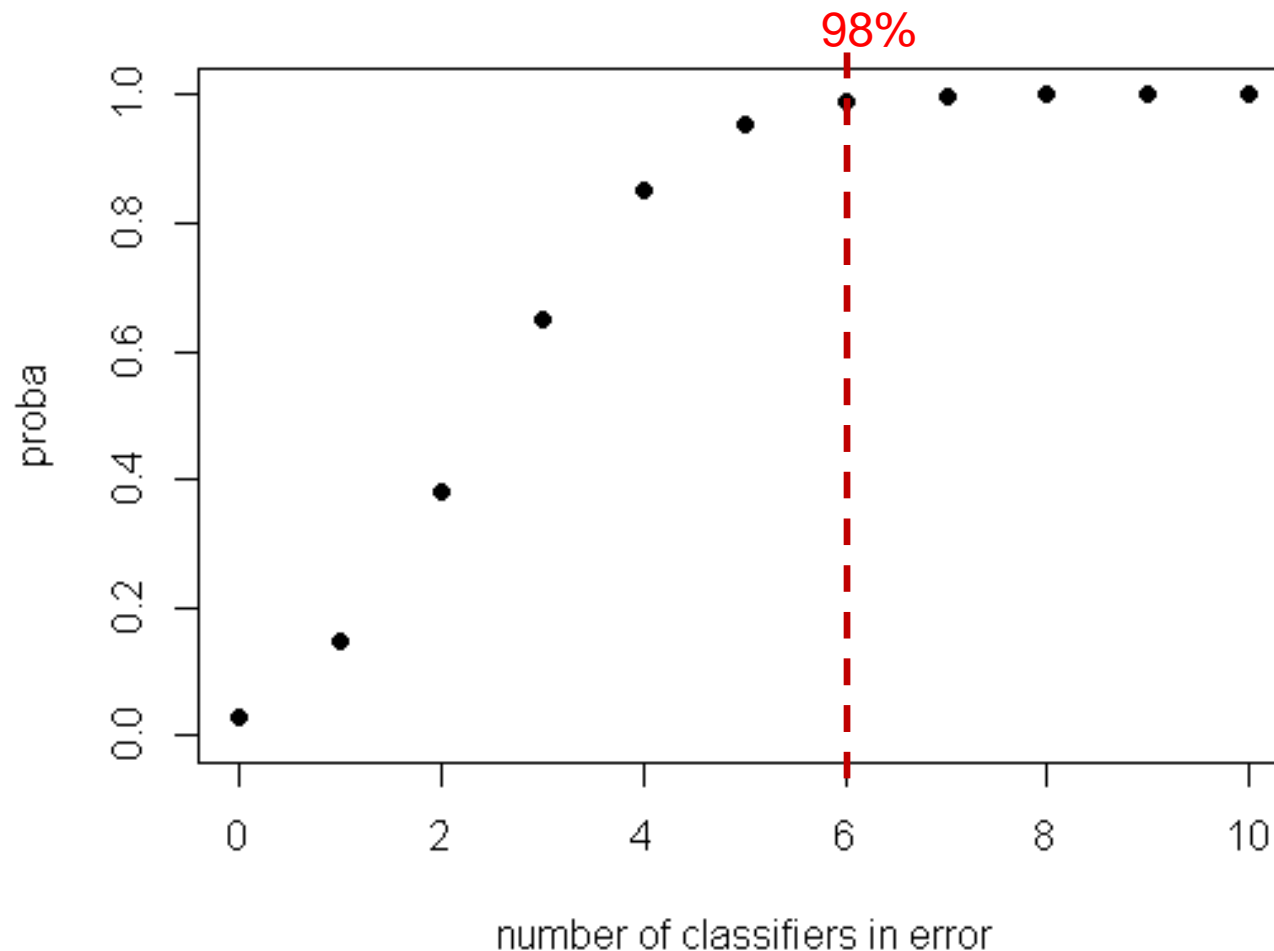
$$p(F(x) \neq y) = 1 - \text{pr}\left(k \leq \frac{M}{2}\right)$$

where pr is the cumulative binomial distribution

The upper bound is much better than the original error rate

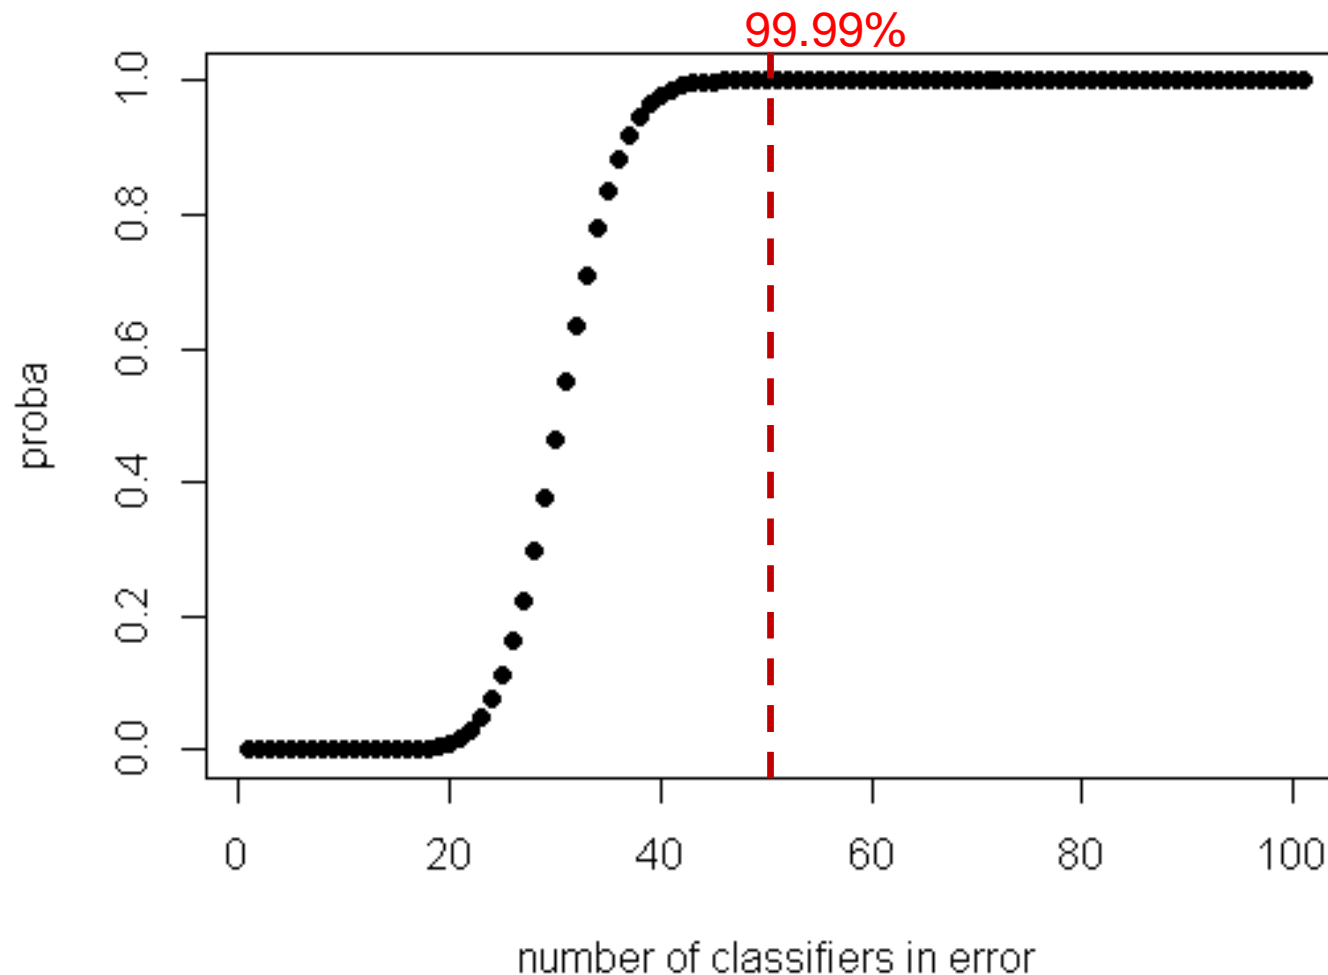
Why to use the ensemble methods

- For 10 classifiers with 30% of error




Why to use the ensemble methods

- For 100 classifiers with 30% of error




Why to use the ensemble methods

10 classifiers




100 classifiers



Number of classifiers with the wrong decision

Why to use the ensemble methods


Statistical reason ↴



From: T. G. Diettrich, Ensemble Methods in Machine Learning, Lecture Notes in Computer Science, Vol. 1857, pages: 1-15, 2000.

Why to use the ensemble methods


Computational reason



From: T. G. Diettrich, Ensemble Methods in Machine Learning, Lecture Notes in Computer Science, Vol. 1857, pages: 1-15, 2000.

Why to use the ensemble methods

Representational reason



From: T. G. Diettrich, Ensemble Methods in Machine Learning, Lecture Notes in Computer Science, Vol. 1857, pages: 1-15, 2000.

How to use ensemble methods

- Set of weak classifiers $p(f(x) \neq y) < 0.5$
 - Train diverse set of models on the same datasets: decision tree, KNN, linear discriminant (generally the simplest, the better).
 - Train different models by using diversity in datasets, parameters, initial conditions.
- Aggregation
 - Bagging
 - Boosting

Bagging

- Create bootstrapped training sets, each containing examples drawn randomly with replacement from the original dataset.
- Train a classifier for each bootstrap dataset
- The final decision is a vote of all classifiers
- Originally developed to reduce the variance of the classifier

Results

Error rate

Data Set	\bar{e}_S	\bar{e}_B	Decrease
waveform	29.0	19.4	33%
heart	10.0	5.3	47%
breast cancer	6.0	4.2	30%
ionosphere	11.2	8.6	23%
diabetes	23.4	18.8	20%
glass	32.0	24.9	22%
soybean	14.5	10.6	27%

Bias-Variance decomposition

- The error rate can be decomposed into a bias and variance term
- Reduction of the variance
- Example : heart dataset


	unaggregated	aggregated
variance	47.64	4.66
bias	1.51	1.97

Adaboost - Adaptive Boosting

- The examples are weighed
 - Each example has a weighth representing its importance in the classification problem.
- AdaBoost is an algorithm constructing complex classifier from a combination of simple and weak classifiers.
- The final decision is based on a weighed vote of weak classifiers.

$$f(x) = \sum_{t=1}^T \alpha_t h_t(x)$$

Toy example




Natural


Not Natural



Toy example




Toy example




Toy example




Toy example




Toy example



Toy example



Toy example




Toy example




Toy example




Toy example



Toy example



Principle



Discrete Adaboost Algorithm

Given: $(x_1, y_1), \dots, (x_m, y_m)$ where $x_i \in X$, $y_i \in Y = \{-1, +1\}$

Initialise $D_1(i) = \frac{1}{m}$.

For $t = 1, \dots, T$:

- Find the classifier $h_t : X \rightarrow \{-1, +1\}$ that minimizes the error with respect to the distribution D_t :

$$h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j, \text{ where } \epsilon_j = \sum_{i=1}^m D_t(i) [y_i \neq h_j(x_i)]$$

- Prerequisite: $\epsilon_t < 0.5$, otherwise stop.
- Choose $\alpha_t \in \mathbf{R}$, typically $\alpha_t = \frac{1}{2} \ln \frac{1 - \epsilon_t}{\epsilon_t}$ where ϵ_t is the weighted error rate of classifier h_t .
- Update:

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

where Z_t is a normalisation factor (chosen so that D_{t+1} will be a distribution).

Output the final classifier:

$$H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$


Find the Weak Classifier

Loop step: Call *WeakLearn*, providing it with the distribution D_t ;
get back weak classifier $h_t : \mathcal{X} \rightarrow \{-1, 1\}$ from $\mathcal{H} = \{h(x)\}$

- ◆ Select a weak classifier with the smallest weighted error
$$h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) [y_i \neq h_j(x_i)]$$
- ◆ Prerequisite: $\epsilon_t < 1/2$ (otherwise stop)
- ◆ *WeakLearn* examples:
 - Decision tree builder, perceptron learning rule – \mathcal{H} infinite
 - Selecting the best one from given *finite* set \mathcal{H}

Demonstration example

Training set



Weak classifier = perceptron

● $\sim N(0, 1)$ ● $\sim \frac{1}{2\pi} e^{-1/2(r-4)^2}$


Find the Weak Classifier

Loop step: Call *WeakLearn*, providing it with the distribution D_t ;
get back weak classifier $h_t : \mathcal{X} \rightarrow \{-1, 1\}$ from $\mathcal{H} = \{h(x)\}$

- ◆ Select a weak classifier with the smallest weighted error
$$h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) [y_i \neq h_j(x_i)]$$
- ◆ Prerequisite: $\epsilon_t < 1/2$ (otherwise stop)
- ◆ *WeakLearn* examples:
 - Decision tree builder, perceptron learning rule – \mathcal{H} infinite
 - Selecting the best one from given *finite* set \mathcal{H}

Demonstration example

Training set



Weak classifier = perceptron

• $\sim N(0, 1)$ • $\sim \frac{1}{2\pi} e^{-1/2(r-4)^2}$

Reweighting

Effect on the training set

Reweighting formula:


$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t} = \frac{\exp(-y_i \sum_{q=1}^t \alpha_q h_q(x_i))}{m \prod_{q=1}^t Z_q}$$

$$\exp(-\alpha_t y_i h_t(x_i)) \begin{cases} < 1, & y_i = h_t(x_i) \\ > 1, & y_i \neq h_t(x_i) \end{cases}$$

$$y * h(x) = 1$$

$$y * h(x) = -1$$

⇒ Increase (decrease) weight of wrongly (correctly) classified examples



Algorithm recapitulation

Initialization...

For $t = 1, \dots, T$:


- ◆ Find $h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) [y_i \neq h_j(x_i)]$
- ◆ If $\epsilon_t \geq 1/2$ then stop
- ◆ Set $\alpha_t = \frac{1}{2} \log\left(\frac{1+r_t}{1-r_t}\right)$
- ◆ Update

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

Output the final classifier:

$$H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$

t = 1



Algorithm recapitulation $t = 1$

Initialization...



For $t = 1, \dots, T$:

- ◆ Find $h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) [y_i \neq h$
- ◆ If $\epsilon_t \geq 1/2$ then stop
- ◆ Set $\alpha_t = \frac{1}{2} \log(\frac{1+r_t}{1-r_t})$
- ◆ Update

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

Output the final classifier:

$$H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$



Algorithm recapitulation

$t = 2$

Initialization...



For $t = 1, \dots, T$:

- ◆ Find $h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) [y_i \neq h_j]$
- ◆ If $\epsilon_t \geq 1/2$ then stop
- ◆ Set $\alpha_t = \frac{1}{2} \log\left(\frac{1+r_t}{1-r_t}\right)$
- ◆ Update

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

Output the final classifier:

$$H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$



Algorithm recapitulation _{$t = 3$}

Initialization...



For $t = 1, \dots, T$:

- ◆ Find $h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) [y_i \neq h_j]$
- ◆ If $\epsilon_t \geq 1/2$ then stop
- ◆ Set $\alpha_t = \frac{1}{2} \log\left(\frac{1+r_t}{1-r_t}\right)$
- ◆ Update

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

Output the final classifier:

$$H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$



Algorithm recap

$t = 4$

Initialization...



For $t = 1, \dots, T$:

- ◆ Find $h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) [y_i \neq h_j]$
- ◆ If $\epsilon_t \geq 1/2$ then stop
- ◆ Set $\alpha_t = \frac{1}{2} \log\left(\frac{1+r_t}{1-r_t}\right)$
- ◆ Update

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

Output the final classifier:

$$H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$




Algorithm recapitulation

$t = 5$

Initialization...

For $t = 1, \dots, T$:


- ◆ Find $h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) [y_i \neq h_j(x_i)]$
- ◆ If $\epsilon_t \geq 1/2$ then stop
- ◆ Set $\alpha_t = \frac{1}{2} \log\left(\frac{1+r_t}{1-r_t}\right)$
- ◆ Update



$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

Output the final classifier:

$$H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$



Algorithm recapitulation $t = 7$

Initialization...



For $t = 1, \dots, T$:

- ◆ Find $h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) [y_i \neq h_j]$
- ◆ If $\epsilon_t \geq 1/2$ then stop
- ◆ Set $\alpha_t = \frac{1}{2} \log\left(\frac{1+r_t}{1-r_t}\right)$
- ◆ Update

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

Output the final classifier:

$$H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$



Algorithm recapitulation $t = 40$

Initialization...



For $t = 1, \dots, T$:

- ◆ Find $h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) [y_i \neq h]$
- ◆ If $\epsilon_t \geq 1/2$ then stop
- ◆ Set $\alpha_t = \frac{1}{2} \log\left(\frac{1+r_t}{1-r_t}\right)$
- ◆ Update

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

Output the final classifier:

$$H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$



Advantages and Drawbacks

Advantages

- Simple to use and implement
- Good performance in generalization

Drawbacks

- Not optimal solution
- Sensitive to noise
- High computing time

Package R

- Ada: `ada()`
- Adabag: `adaboost.M1()`
- Bst: `bst()`
- Mboost: `glmboost()`
- wSVM: `wsvm()`